



Tecnológico de Monterrey

Proyecto BD Clasificador de Imágenes

Alejandro Delgado Medrano A01227074

Bases de Datos Avanzadas

GRUPO 01

Profesor Rodolfo Rubén Álvarez González

En el presente proyecto se hará una comparación de imágenes, usando tensorflow para enseñarle a la computadora a distinguirlas a base de carpetas de las mismas. En este caso son 3 carpetas de perros de diferentes razas, al juntar un numero mayor a 50 imágenes en cada carpeta el programa usando machine learning, aprenderá a diferenciarlos. Al igual para probar que si sea cierto se agregaran imágenes nuevas en otra carpeta en la cual el programa las revisara y comparara mostrando al final del proceso cual raza son y cual es su porcentaje de seguridad en ello.

Tensorflow:

Es una librería de código libre para Machine Learning.

TensorFlow fue desarrollado originalmente por investigadores e ingenieros que trabajaban en el equipo de Google Brain Team, dentro del departamento de investigación de Machine Intelligence, con el propósito de llevar a cabo el aprendizaje automático y la investigación de redes neuronales profundas.

TensorFlow es una herramienta de machine learning, popularizada por su eficiencia con redes neuronales de aprendizaje profundo pero que permite la ejecución de procesos distribuidos que no tengan nada que ver con redes neuronales.

Tiene múltiples usos, por ejemplo:

- Puede **reconocer varias palabras** del alfabeto porque relaciona las letras y fonemas.
- **imágenes y textos** que se pueden relacionar entre sí rápidamente gracias a la capacidad de asociación del sistema de redes neuronales.
- Motor de reconocimiento de imágenes **DeepDream**
- mejorar la fotografía de los smartphones
- Para ayudar al diagnóstico médico

La arquitectura flexible de TensorFlow le permite implementar el cálculo a una o más CPU o GPU en equipos de escritorio, servidores o dispositivos móviles con una sola API. Tiene muchas arquitecturas, pero yo use el driver de Spark como gestor de parámetros

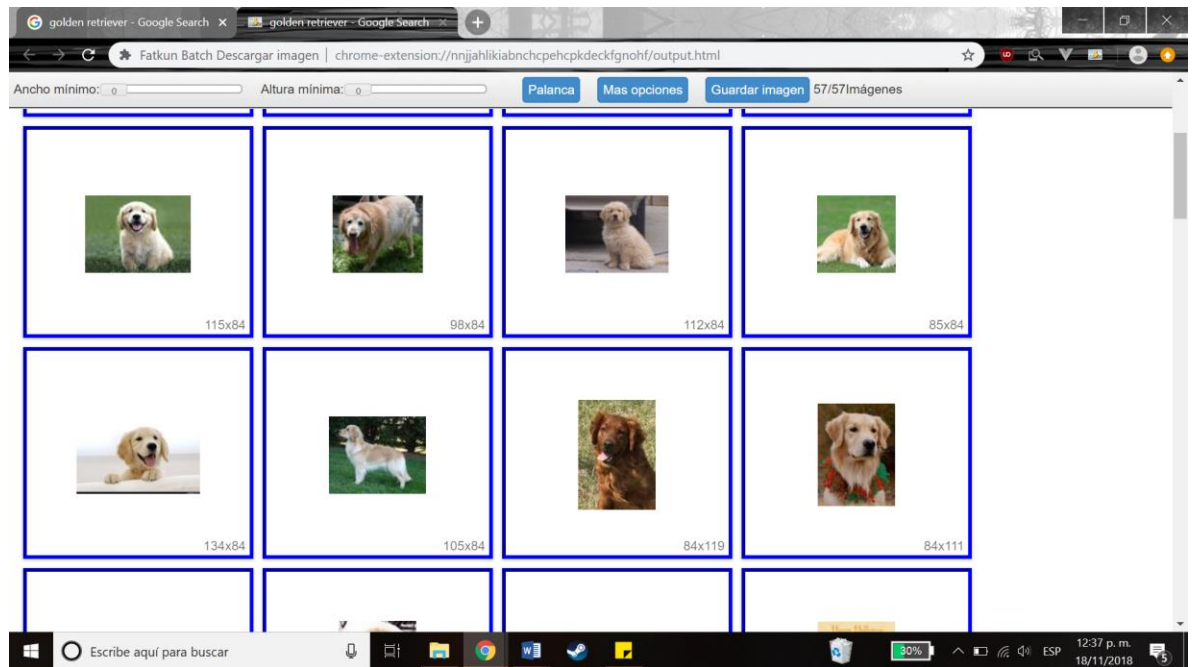
Ventajas:

- Tiene flexibilidad de plataforma.
- Es fácil de entrenar en CPU y GPU para computación distribuida.
- TensorFlow tiene capacidades de diferenciación automática
- Tiene soporte avanzado para subprocesos, cómputo asíncrono y colas.

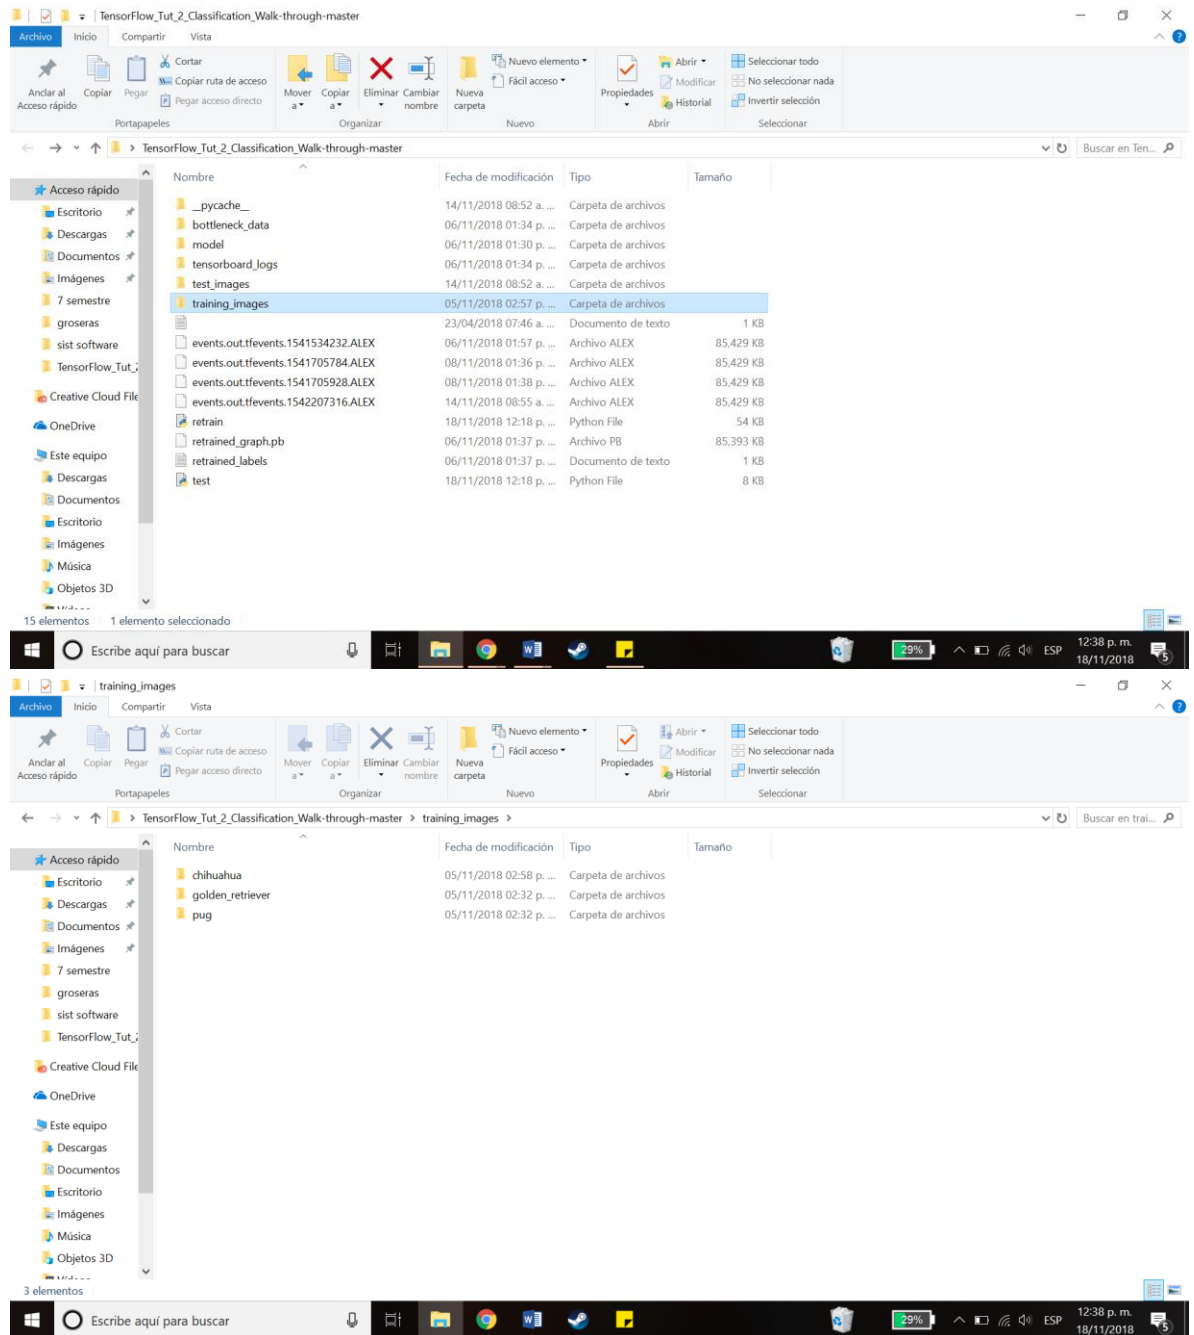
- Es un código personalizable y de código abierto.

Demostración:

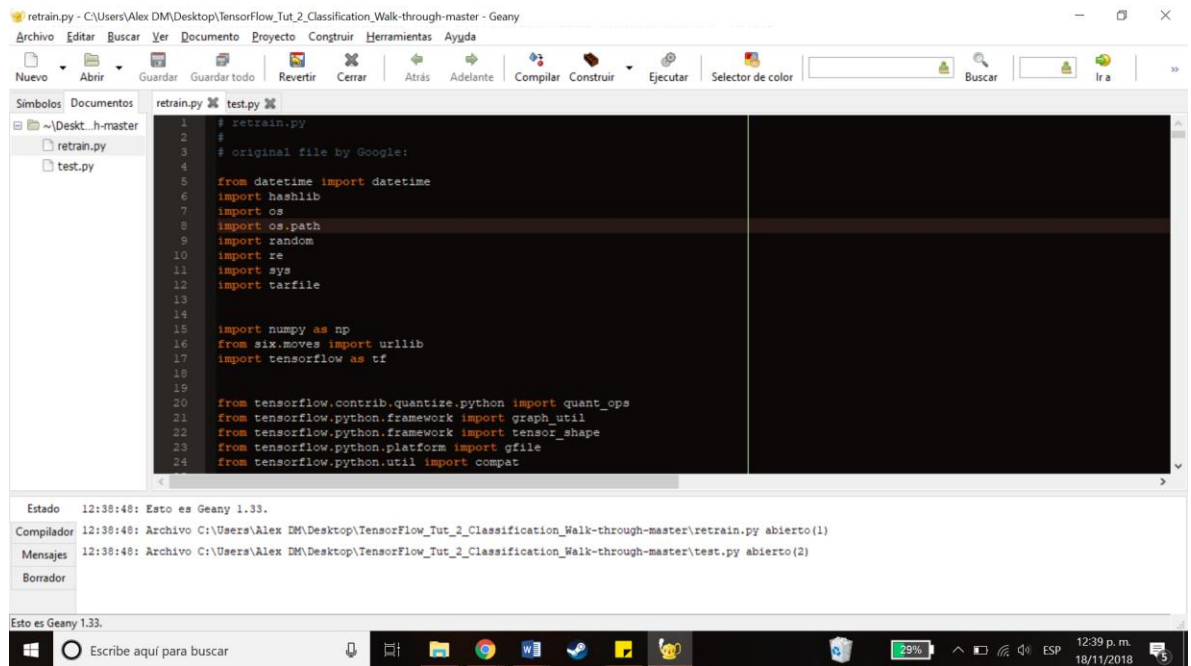
Lo primero que se tiene que hacer es la recolección de imágenes, para esto recomiendo usar una herramienta de Google Chrom llamado “Fatkun” con la cual se selecciona una categoría de imágenes a base de Google y las descarga en la máquina para su uso.




Después de eso se deben agregar todas esas imágenes en la carpeta del proyecto llamado “training_images”

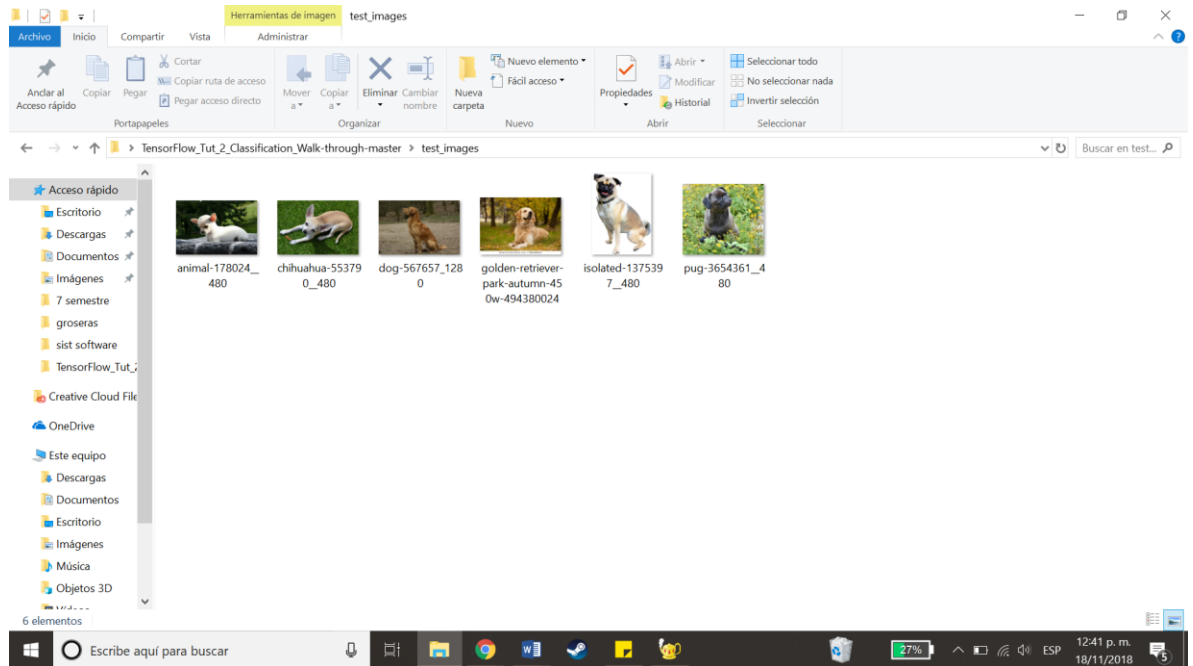


Para el siguiente paso abrimos el código de “retrain” (yo uso Geany por su simpleza y efectividad) el cual es el que se encarga de hacer el machine learning con las imágenes ya descargadas (el proceso puede tardar un rato)

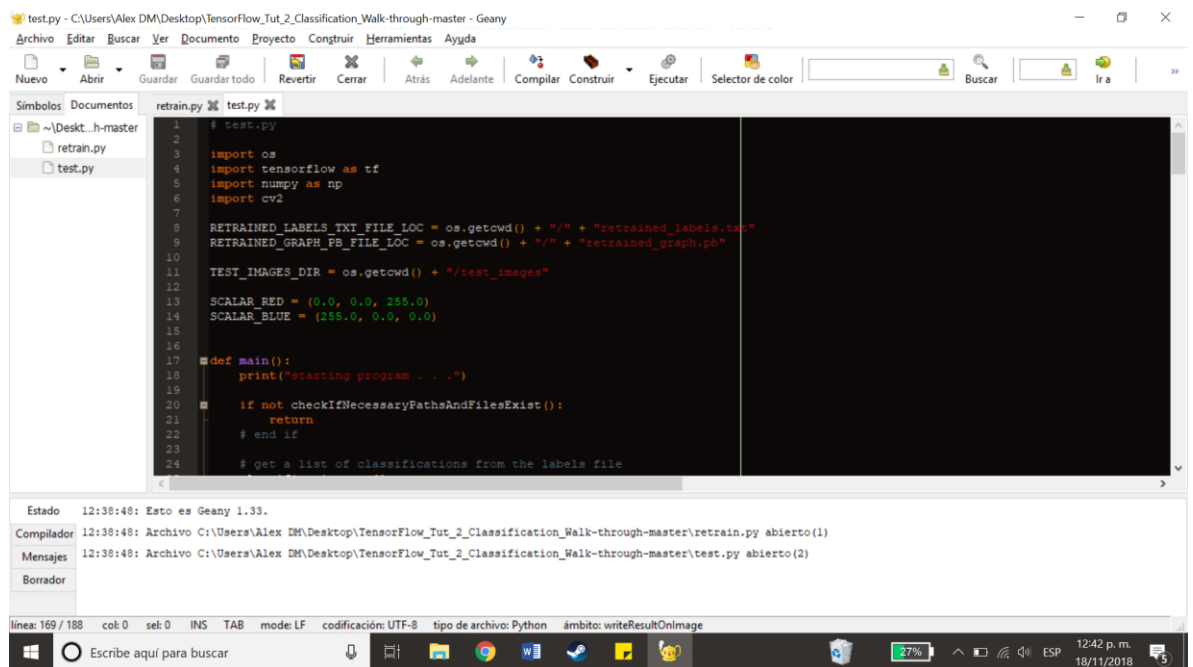


Una vez terminado de correr elegimos nuevas imágenes de las categorías ya descargadas (con 1 o 2 es suficiente) y las insertamos en la carpeta de “test_images”

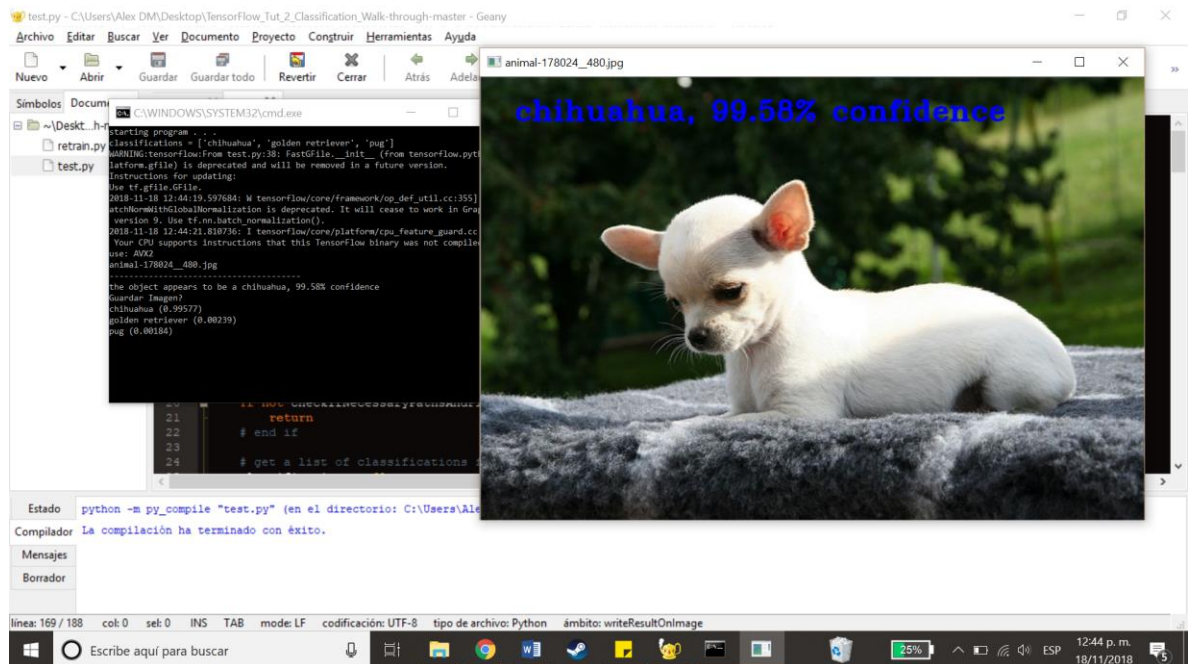




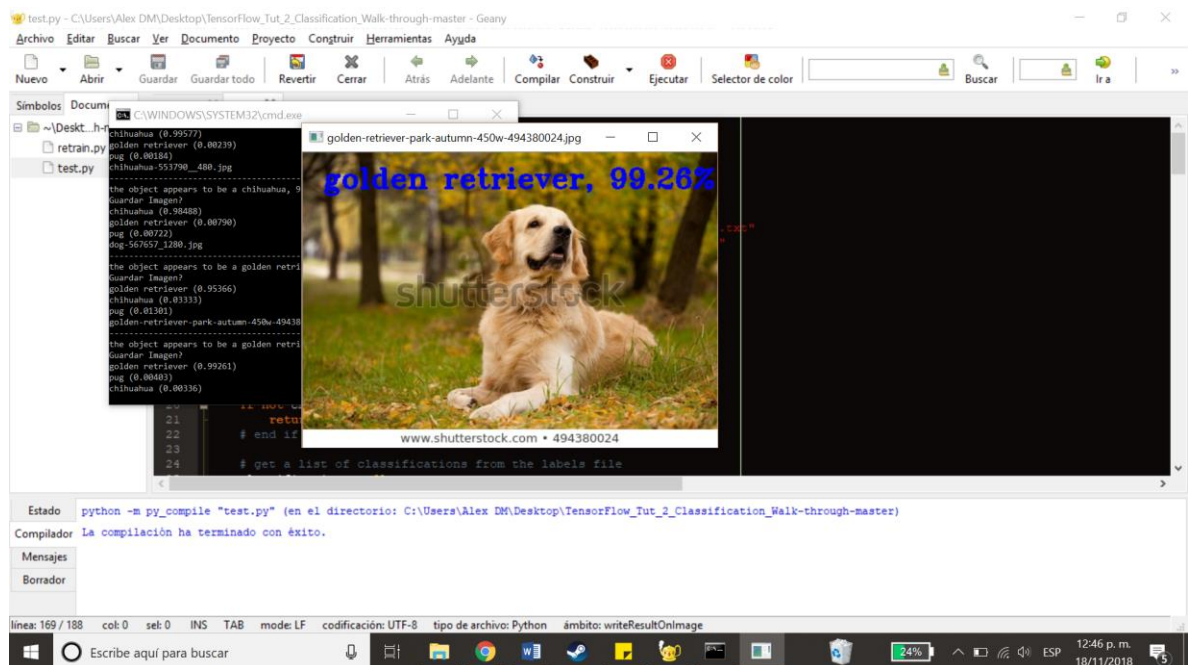
Una vez terminado abrimos el código de “test” en el cual se correra lo mas importante del proyecto para probar que la base de datos si sirva, al igual que la maquina si pueda diferenciar las imágenes.



Al terminar de correr el código se mostrara en pantalla los resultados de las imágenes:



como podemos ver la primera imagen la identifico como un Chihuahua con una seguridad de 99.58%



Al igual nos muestra lo mismo de la siguiente, pero con un poco menos de seguridad.

Una vez terminado el programa y las imágenes en esa carpeta el CMD se cerrara y esta terminado el ejemplo.

Código:

En el presente código se explicaran la función de cada parte de código, usando los comentarios y la explicación de porque fue hecho así.

En el primero tenemos el código de “retrain”, en este código es donde se le enseña a la base de datos a base de Machine learning a aprender a diferenciar las imágenes basándonos en las carpetas establecidas y lograr un rango del 95% al 99% de seguridad.

```
# retrain.py

# original file by Google:

from datetime import datetime

import hashlib
import os
import os.path
import random
import re
import sys
import tarfile

import numpy as np

from six.moves import urllib

import tensorflow as tf

from tensorflow.contrib.quantize.python import
quant_ops

from tensorflow.python.framework import
graph_util

from tensorflow.python.framework import
tensor_shape

from tensorflow.python.platform import gfile

from tensorflow.python.util import compat

# module level variables

MIN_NUM_IMAGES_REQUIRED_FOR_TRAINING =
10

MIN_NUM_IMAGES_SUGGESTED_FOR_TRAINING
= 100

MIN_NUM_IMAGES_REQUIRED_FOR_TESTING = 3

MAX_NUM_IMAGES_PER_CLASS = 2 ** 27 - 1 #
~134M

# path to folders of labeled images

TRAINING_IMAGES_DIR = os.getcwd() +
'/training_images'

TEST_IMAGES_DIR = os.getcwd() + "/test_images/"

# where to save the trained graph

OUTPUT_GRAPH = os.getcwd() + '/' +
'retrained_graph.pb'

# where to save the intermediate graphs

INTERMEDIATE_OUTPUT_GRAPHS_DIR =
os.getcwd() + '/intermediate_graph'

# how many steps to store intermediate graph, if
"0" then will not store

INTERMEDIATE_STORE_FREQUENCY = 0

# where to save the trained graph's labels

OUTPUT_LABELS = os.getcwd() + '/' +
'retrained_labels.txt'

# where to save summary logs for TensorBoard

TENSORBOARD_DIR = os.getcwd() + '/' +
'tensorboard_logs'

# how many training steps to run before ending
```


NOTE: original Google default is 4000, use 4000
(or possibly higher) for production grade results

HOW_MANY_TRAINING_STEPS=500

how large a learning rate to use when training

LEARNING_RATE = 0.01

what percentage of images to use as a test set

TESTING_PERCENTAGE = 10

what percentage of images to use as a validation
set

VALIDATION_PERCENTAGE = 10

how often to evaluate the training results

EVAL_STEP_INTERVAL = 10

how many images to train on at a time

TRAIN_BATCH_SIZE = 100

How many images to test on. This test set is only
used once, to evaluate the final accuracy of the
model after

training completes. A value of -1 causes the
entire test set to be used, which leads to more
stable results across runs.

TEST_BATCH_SIZE = -1

How many images to use in an evaluation batch.
This validation set is used much more often than
the test set, and is an early indicator of how

accurate the model is during training. A value of -
1 causes the entire validation set to be used, which
leads to

more stable results across training iterations, but
may be slower on large training sets.

VALIDATION_BATCH_SIZE = 100

whether to print out a list of all misclassified test
images

PRINT_MISCLASSIFIED_TEST_IMAGES = False

Path to classify_image_graph_def.pb,
imagenet_synset_to_human_label_map.txt, and
imagenet_2012_challenge_label_map_proto.pbtxt

MODEL_DIR = os.getcwd() + "/" + "model"

Path to cache bottleneck layer values as files

BOTTLENECK_DIR = os.getcwd() + '/' +
'bottleneck_data'

the name of the output classification layer in the
retrained graph

FINAL_TENSOR_NAME = 'final_result'

whether to randomly flip half of the training
images horizontally

FLIP_LEFT_RIGHT = False

a percentage determining how much of a margin
to randomly crop off the training images

RANDOM_CROP = 0

a percentage determining how much to
randomly scale up the size of the training images
by

RANDOM_SCALE = 0

a percentage determining how much to
randomly multiply the training image input pixels
up or down by

RANDOM_BRIGHTNESS = 0

ARCHITECTURE = 'inception_v3'

def main():

```

print("starting program . . .")

# make sure the logging output is visible, see
https://github.com/tensorflow/tensorflow/issues/3047

tf.logging.set_verbosity(tf.logging.INFO)

if not checkIfNecessaryPathsAndFilesExist():

    return

# end if

# prepare necessary directories that can be used
during training

prepare_file_system()

# Gather information about the model
architecture we'll be using.

model_info =
create_model_info(ARCHITECTURE)

if not model_info:

    tf.logging.error('Did not recognize architecture
flag')

    return -1

# end if

# download the model if necessary, then create
the model graph

print("downloading model (if necessary) . . .")

downloadModelIfNotAlreadyPresent(model_info['
data_url'])

print("creating model graph . . .")

graph, bottleneck_tensor, resized_image_tensor
= (create_model_graph(model_info))

# Look at the folder structure, and create lists of
all the images.

print("creating image lists . . .")

image_lists =
create_image_lists(TRAINING_IMAGES_DIR,

```

```

TESTING_PERCENTAGE,
VALIDATION_PERCENTAGE)

class_count = len(image_lists.keys())

if class_count == 0:

    tf.logging.error('No valid folders of images
found at ' + TRAINING_IMAGES_DIR)

    return -1

# end if

if class_count == 1:

    tf.logging.error('Only one valid folder of
images found at ' + TRAINING_IMAGES_DIR + ' -
multiple classes are needed for classification.')

    return -1

# end if

# determine if any of the distortion command
line flags have been set

doDistortImages = False

if (FLIP_LEFT_RIGHT == True or RANDOM_CROP
!= 0 or RANDOM_SCALE != 0 or
RANDOM_BRIGHTNESS != 0):

    doDistortImages = True

# end if

print("starting session . . .")

with tf.Session(graph=graph) as sess:

    # Set up the image decoding sub-graph.

    print("performing jpeg decoding . . .")

    jpeg_data_tensor, decoded_image_tensor =
add_jpeg_decoding( model_info['input_width'],

model_info['input_height'],

model_info['input_depth'],

model_info['input_mean'],

model_info['input_std'])

    print("caching bottlenecks . . .")

```

```

        distorted_jpeg_data_tensor = None

        distorted_image_tensor = None

        if doDistortImages:

            # We will be applying distortions, so setup
            the operations we'll need.

            (distorted_jpeg_data_tensor,
            distorted_image_tensor) =
            add_input_distortions(FLIP_LEFT_RIGHT,
            RANDOM_CROP, RANDOM_SCALE,

            RANDOM_BRIGHTNESS,
            model_info['input_width'],

            model_info['input_height'],
            model_info['input_depth'],

            model_info['input_mean'],
            model_info['input_std'])

        else:

            # We'll make sure we've calculated the
            'bottleneck' image summaries and

            # cached them on disk.

            cache_bottlenecks(sess, image_lists,
            TRAINING_IMAGES_DIR, BOTTLENECK_DIR,
            jpeg_data_tensor, decoded_image_tensor,

            resized_image_tensor,
            bottleneck_tensor, ARCHITECTURE)

            # end if

            # Add the new layer that we'll be training.

            print("adding final training layer . . .")

            (train_step, cross_entropy, bottleneck_input,
            ground_truth_input, final_tensor) =
            add_final_training_ops(len(image_lists.keys()),

            FINAL_TENSOR_NAME,

            bottleneck_tensor,

            model_info['bottleneck_tensor_size'],

            model_info['quantize_layer'])

            # Create the operations we need to evaluate
            the accuracy of our new layer.

            print("adding eval ops for final training layer . .
            .")

            evaluation_step, prediction =
            add_evaluation_step(final_tensor,
            ground_truth_input)

            # Merge all the summaries and write them out
            to the tensorboard_dir

            print("writing TensorBoard info . . .")

            merged = tf.summary.merge_all()

            train_writer =
            tf.summary.FileWriter(TENSORBOARD_DIR +
            '/train', sess.graph)

            validation_writer =
            tf.summary.FileWriter(TENSORBOARD_DIR +
            '/validation')

            # Set up all our weights to their initial default
            values.

            init = tf.global_variables_initializer()

            sess.run(init)

            # Run the training for as many cycles as
            requested on the command line.

            print("performing training . . .")

            for i in range(HOW_MANY_TRAINING_STEPS):

                # Get a batch of input bottleneck values,
                either calculated fresh every

                # time with distortions applied, or from the
                cache stored on disk.

                if doDistortImages:

                    (train_bottlenecks, train_ground_truth) =
                    get_random_distorted_bottlenecks(sess,
                    image_lists, TRAIN_BATCH_SIZE, 'training',

                    TRAINING_IMAGES_DIR,
                    distorted_jpeg_data_tensor,

                    distorted_image_tensor, resized_image_tensor,
                    bottleneck_tensor)

                else:

```

```

        (train_bottlenecks, train_ground_truth, _)
    = get_random_cached_bottlenecks(sess,
    image_lists, TRAIN_BATCH_SIZE, 'training',

```

```

BOTTLENECK_DIR, TRAINING_IMAGES_DIR,
jpeg_data_tensor,

```

```

decoded_image_tensor, resized_image_tensor,
bottleneck_tensor,

```

```

ARCHITECTURE)

```

```

    # end if

```

```

        # Feed the bottlenecks and ground truth
        into the graph, and run a training

```

```

        # step. Capture training summaries for
        TensorBoard with the `merged` op.

```

```

        train_summary, _ = sess.run([merged,
        train_step], feed_dict={bottleneck_input:
        train_bottlenecks, ground_truth_input:
        train_ground_truth})

```

```

        train_writer.add_summary(train_summary,
        i)

```

```

        # Every so often, print out how well the
        graph is training.

```

```

        is_last_step = (i + 1 ==
        HOW_MANY_TRAINING_STEPS)

```

```

        if (i % EVAL_STEP_INTERVAL) == 0 or
        is_last_step:

```

```

            train_accuracy, cross_entropy_value =
            sess.run([evaluation_step, cross_entropy],
            feed_dict={bottleneck_input: train_bottlenecks,
            ground_truth_input: train_ground_truth})

```

```

            tf.logging.info('%s: Step %d: Train
            accuracy = %.1f%%' % (datetime.now(), i,
            train_accuracy * 100))

```

```

            tf.logging.info('%s: Step %d: Cross entropy
            = %f' % (datetime.now(), i, cross_entropy_value))

```

```

            validation_bottlenecks,
            validation_ground_truth, _ =
            (get_random_cached_bottlenecks(sess,
            image_lists, VALIDATION_BATCH_SIZE, 'validation',

```

```

BOTTLENECK_DIR, TRAINING_IMAGES_DIR,
jpeg_data_tensor,

```

```

decoded_image_tensor, resized_image_tensor,
bottleneck_tensor,

```

```

ARCHITECTURE))

```

```

        # Run a validation step and capture
        training summaries for TensorBoard with the
        `merged` op.

```

```

        validation_summary, validation_accuracy
    = sess.run(

```

```

        [merged, evaluation_step],
        feed_dict={bottleneck_input:
        validation_bottlenecks, ground_truth_input:
        validation_ground_truth})

```

```

        validation_writer.add_summary(validation_summa
        ry, i)

```

```

        tf.logging.info('%s: Step %d: Validation
        accuracy = %.1f%% (N=%d)' % (datetime.now(), i,
        validation_accuracy * 100,
        len(validation_bottlenecks)))

```

```

    # end if

```

```

        # Store intermediate results

```

```

        intermediate_frequency =
        INTERMEDIATE_STORE_FREQUENCY

```

```

        if (intermediate_frequency > 0 and (i %
        intermediate_frequency == 0) and i > 0):

```

```

            intermediate_file_name =
            (INTERMEDIATE_OUTPUT_GRAPHS_DIR +
            'intermediate_' + str(i) + '.pb')

```

```

            tf.logging.info('Save intermediate result to
            : ' + intermediate_file_name)

```

```

            save_graph_to_file(sess, graph,
            intermediate_file_name)

```

```

        # end if

```

```

    # end for

```

```

        # We've completed all our training, so run a
        final test evaluation on some new images we
        haven't used before

```

```

        print("running testing . . .")

```

```

        test_bottlenecks, test_ground_truth,
        test_filenames =

```

```
(get_random_cached_bottlenecks(sess,
image_lists, TEST_BATCH_SIZE, 'testing',
BOTTLENECK_DIR,
```

```
TRAINING_IMAGES_DIR, jpeg_data_tensor,
decoded_image_tensor, resized_image_tensor,
```

```
bottleneck_tensor, ARCHITECTURE))
```

```
test_accuracy, predictions =
sess.run([evaluation_step, prediction],
feed_dict={bottleneck_input: test_bottlenecks,
ground_truth_input: test_ground_truth})
```

```
tf.logging.info('Final test accuracy = %.1f%%
(N=%d)' % (test_accuracy * 100,
len(test_bottlenecks)))
```

```
if PRINT_MISCLASSIFIED_TEST_IMAGES:
```

```
tf.logging.info('=== MISCLASSIFIED TEST
IMAGES ===')
```

```
for i, test_filename in
enumerate(test_filenames):
```

```
if predictions[i] != test_ground_truth[i]:
```

```
tf.logging.info('%70s %s' %
(test_filename,
list(image_lists.keys())[predictions[i]]))
```

```
# end if
```

```
# end for
```

```
# end if
```

```
# write out the trained graph and labels with
the weights stored as constants
```

```
print("writing trained graph and labbels with
weights")
```

```
save_graph_to_file(sess, graph,
OUTPUT_GRAPH)
```

```
with gfile.FastGFile(OUTPUT_LABELS, 'w') as f:
```

```
f.write('\n'.join(image_lists.keys()) + '\n')
```

```
# end with
```

```
print("done !!!")
```

```
def checkIfNecessaryPathsAndFilesExist():
```

```
# if the training directory does not exist, show
and error message and bail
```

```
if not os.path.exists(TRAINING_IMAGES_DIR):
```

```
print("")
```

```
print('ERROR: TRAINING_IMAGES_DIR "" +
TRAINING_IMAGES_DIR + "" does not seem to
exist')
```

```
print('Did you set up the training images?')
```

```
print("")
```

```
return False
```

```
# end if
```

```
# nested class
```

```
class TrainingSubDir:
```

```
# constructor
```

```
def __init__(self):
```

```
self.loc = ""
```

```
self.numImages = 0
```

```
# end constructor
```

```
# end class
```

```
# declare a list of training sub-directories
```

```
trainingSubDirs = []
```

```
# populate the training sub-directories
```

```
for dirName in
os.listdir(TRAINING_IMAGES_DIR):
```

```
currentTrainingImagesSubDir =
os.path.join(TRAINING_IMAGES_DIR, dirName)
```

```
if os.path.isdir(currentTrainingImagesSubDir):
```

```
trainingSubDir = TrainingSubDir()
```

```
trainingSubDir.loc =
currentTrainingImagesSubDir
```

```
trainingSubDirs.append(trainingSubDir)
```

```
# end if
```

```
# end for
```

```
# if no training sub-directories were found, show
an error message and return false
```

```
if len(trainingSubDirs) == 0:
```

```
    print("ERROR: there don't seem to be any
training image sub-directories in " +
TRAINING_IMAGES_DIR)
```

```
    print("Did you make a separate image sub-
directory for each classification type?")
```

```
    return False
```

```
# end if
```

```
# populate the number of training images in
each training sub-directory
```

```
for trainingSubDir in trainingSubDirs:
```

```
    # count how many images are in the current
training sub-directory
```

```
    for fileName in os.listdir(trainingSubDir.loc):
```

```
        if fileName.endswith(".jpg"):
```

```
            trainingSubDir.numImages += 1
```

```
# if any training sub-directory has less than the
min required number of training images, show an
error message and return false
```

```
for trainingSubDir in trainingSubDirs:
```

```
    if trainingSubDir.numImages <
MIN_NUM_IMAGES_REQUIRED_FOR_TRAINING:
```

```
        print("ERROR: there are less than the
required " +
str(MIN_NUM_IMAGES_REQUIRED_FOR_TRAINING) + "
images in " + trainingSubDir.loc)
```

```
        print("Did you populate each training sub-
directory with images?")
```

```
        return False
```

```
    # end if
```

```
# end for
```

```
# if any training sub-directory has less than the
recommended number of training images, show a
warning (but don't return false)
```

```
for trainingSubDir in trainingSubDirs:
```

```
    if trainingSubDir.numImages <
MIN_NUM_IMAGES_SUGGESTED_FOR_TRAINING:
```

```
        print("WARNING: there are less than the
suggested " +
str(MIN_NUM_IMAGES_SUGGESTED_FOR_TRAINING) + "
images in " + trainingSubDir.loc)
```

```
        print("More images should be added to this
directory for acceptable training results")
```

```
        # note we do not return false here b/c this is
a warning, not an error
```

```
# if the test images directory does not exist,
show an error message and bail
```

```
if not os.path.exists(TEST_IMAGES_DIR):
```

```
    print("")
```

```
    print('ERROR: TEST_IMAGES_DIR "" +
TEST_IMAGES_DIR + "" does not seem to exist')
```

```
    print('Did you break out some test images?')
```

```
    print("")
```

```
    return False
```

```
# end if
```

```
# count how many images are in the test images
directory
```

```
numImagesInTestDir = 0
```

```
for fileName in os.listdir(TEST_IMAGES_DIR):
```

```
    if fileName.endswith(".jpg"):
```

```
        numImagesInTestDir += 1
```

```
# if there are not enough images in the test
images directory, show an error and return false
```

```
if numImagesInTestDir <
MIN_NUM_IMAGES_REQUIRED_FOR_TESTING:
```

```
    print("ERROR: there are not at least " +
str(MIN_NUM_IMAGES_REQUIRED_FOR_TESTING) + "
images in " + TEST_IMAGES_DIR)
```

```

        print("Did you break out some test images?")

    return False

return True

```

```
def prepare_file_system():
```

```

    # Setup the directory we'll write summaries to
    for TensorBoard

```

```
    if tf.gfile.Exists(TENSORBOARD_DIR):
```

```
        tf.gfile.DeleteRecursively(TENSORBOARD_DIR)
```

```
    tf.gfile.MakeDirs(TENSORBOARD_DIR)
```

```
    if INTERMEDIATE_STORE_FREQUENCY > 0:
```

```

    makeDirIfDoesNotExist(INTERMEDIATE_OUTPUT_
    GRAPHS_DIR)

```

```
    return
```

```
def makeDirIfDoesNotExist(dir_name):
```

```
    """
```

```
    Makes sure the folder exists on disk.
```

```
    Args:
```

```

        dir_name: Path string to the folder we want to
        create.
    """

```

```
    """
```

```
    if not os.path.exists(dir_name):
```

```
        os.makedirs(dir_name)
```

```
def create_model_info(architecture):
```

```
    """
```

```

    Given the name of a model architecture, returns
    information about it.

```

```

    There are different base image recognition
    pretrained models that can be

```

```

    retrained using transfer learning, and this
    function translates from the name

```

```

    of a model to the attributes that are needed to
    download and train with it.

```

Args:

architecture: Name of a model architecture.

Returns:

Dictionary of information about the model, or
None if the name isn't recognized

Raises:

ValueError: If architecture name is unknown.

```
    """
```

```
    architecture = architecture.lower()
```

```
    is_quantized = False
```

```
    if architecture == 'inception_v3':
```

```
        # pylint: disable=line-too-long
```

```

        data_url =
        'http://download.tensorflow.org/models/image/i
        magenet/inception-2015-12-05.tgz'

```

```
        # pylint: enable=line-too-long
```

```

        bottleneck_tensor_name =
        'pool_3/_reshape:0'

```

```
        bottleneck_tensor_size = 2048
```

```
        input_width = 299
```

```
        input_height = 299
```

```
        input_depth = 3
```

```
        resized_input_tensor_name = 'Mul:0'
```

```

        model_file_name =
        'classify_image_graph_def.pb'

```

```
        input_mean = 128
```

```
        input_std = 128
```

```
    elif architecture.startswith('mobilenet_'):
```

```
        parts = architecture.split('_')
```

```
        if len(parts) != 3 and len(parts) != 4:
```

```

            tf.logging.error("Couldn't understand
            architecture name '%s'", architecture)

```

```
    return None
```

```

# end if

version_string = parts[1]

if (version_string != '1.0' and version_string !=
'0.75' and version_string != '0.50' and
version_string != '0.25'):

    tf.logging.error("""The Mobilenet version
should be '1.0', '0.75', '0.50', or '0.25', but found
'%s' for architecture '%s'""" % (version_string,
architecture))

    return None

# end if

size_string = parts[2]

if (size_string != '224' and size_string != '192'
and size_string != '160' and size_string != '128'):

    tf.logging.error("""The Mobilenet input size
should be '224', '192', '160', or '128', but found '%s'
for architecture '%s'""" % (size_string, architecture))

    return None

if len(parts) == 3:

    is_quantized = False

else:

    if parts[3] != 'quantized':

        tf.logging.error(

            "Couldn't understand architecture suffix
'%s' for '%s'" % (parts[3], architecture))

        return None

    is_quantized = True

if is_quantized:

    data_url =
'http://download.tensorflow.org/models/mobilenet
_v1_'

    data_url += version_string + '_' + size_string
+ '_quantized_frozen.tgz'

    bottleneck_tensor_name =
'MobilenetV1/Predictions/Reshape:0'

    resized_input_tensor_name =
'Placeholder:0'

```

```

    model_dir_name = ('mobilenet_v1_' +
version_string + '_' + size_string +
'_quantized_frozen')

    model_base_name =
'quantized_frozen_graph.pb'

    else:

        data_url =
'http://download.tensorflow.org/models/mobilenet
_v1_'

        data_url += version_string + '_' + size_string
+ '_frozen.tgz'

        bottleneck_tensor_name =
'MobilenetV1/Predictions/Reshape:0'

        resized_input_tensor_name = 'input:0'

        model_dir_name = ('mobilenet_v1_' +
version_string + '_' + size_string

        model_base_name = 'frozen_graph.pb'

        bottleneck_tensor_size = 1001

        input_width = int(size_string)

        input_height = int(size_string)

        input_depth = 3

        model_file_name =
os.path.join(model_dir_name, model_base_name)

        input_mean = 127.5

        input_std = 127.5

    else:

        tf.logging.error("Couldn't understand
architecture name '%s'", architecture)

        raise ValueError('Unknown architecture',
architecture)

    return {'data_url': data_url,
'bottleneck_tensor_name':
bottleneck_tensor_name,
'bottleneck_tensor_size': bottleneck_tensor_size,

        'input_width': input_width, 'input_height':
input_height, 'input_depth': input_depth,
'resized_input_tensor_name':
resized_input_tensor_name,

        'model_file_name': model_file_name,
'input_mean': input_mean, 'input_std': input_std,
'quantize_layer': is_quantized, }

```



```
def downloadModelIfNotAlreadyPresent(data_url):
```

```
    """
```

```
    Download and extract model tar file.
```

```
    If the pretrained model we're using doesn't
    already exist, this function downloads it from the
    TensorFlow.org website and unpacks it into a
    directory.
```

```
    Args:
```

```
        data_url: Web location of the tar file
        containing the pretrained model.
```

```
    """
```

```
    dest_directory = MODEL_DIR
```

```
    if not os.path.exists(dest_directory):
```

```
        os.makedirs(dest_directory)
```

```
    # end if
```

```
    filename = data_url.split('/')[-1]
```

```
    filepath = os.path.join(dest_directory, filename)
```

```
    if not os.path.exists(filepath):
```

```
        # nested function
```

```
        def _progress(count, block_size, total_size):
```

```
            sys.stdout.write('\r>> Downloading %s
            %.1f%%' % (filename, float(count * block_size) /
            float(total_size) * 100.0))
```

```
            sys.stdout.flush()
```

```
        # end def
```

```
    filepath, _ = urllib.request.urlretrieve(data_url,
    filepath, _progress)
```

```
    print()
```

```
    statinfo = os.stat(filepath)
```

```
    tf.logging.info('Successfully downloaded ' +
    str(filename) + ', statinfo.st_size = ' +
    str(statinfo.st_size) + ' bytes')
```

```
    print('Extracting file from ', filepath)
```

```
        tarfile.open(filepath,
        'r:gz').extractall(dest_directory)
```

```
    else:
```

```
        print('Not extracting or downloading files,
        model already present in disk')
```

```
def create_model_graph(model_info):
```

```
    """
```

```
    Creates a graph from saved GraphDef file and
    returns a Graph object.
```

```
    Args:
```

```
        model_info: Dictionary containing information
        about the model architecture.
```

```
    Returns:
```

```
        Graph holding the trained Inception network,
        and various tensors we'll be manipulating.
```

```
    """
```

```
    with tf.Graph().as_default() as graph:
```

```
        model_path = os.path.join(MODEL_DIR,
        model_info['model_file_name'])
```

```
        print('Model path: ', model_path)
```

```
        with gfile.FastGFile(model_path, 'rb') as f:
```

```
            graph_def = tf.GraphDef()
```

```
            graph_def.ParseFromString(f.read())
```

```
            bottleneck_tensor, resized_input_tensor =
            (tf.import_graph_def(graph_def, name="",
            return_elements=[model_info['bottleneck_tensor_
            name'],
            model_info['resized_input_tensor_name'],]))
```

```
        return graph, bottleneck_tensor,
        resized_input_tensor
```

```
def create_image_lists(image_dir,
    testing_percentage, validation_percentage):
```

```
    """
```

```
    Builds a list of training images from the file
    system.
```

Analyzes the sub folders in the image directory, splits them into stable

training, testing, and validation sets, and returns a data structure

describing the lists of images for each label and their paths.

Args:

image_dir: String path to a folder containing subfolders of images.

testing_percentage: Integer percentage of the images to reserve for tests.

validation_percentage: Integer percentage of images reserved for validation.

Returns:

A dictionary containing an entry for each label subfolder, with images split

into training, testing, and validation sets within each label.

"""

if the image directory does not exist, log an error and bail

if not gfile.Exists(image_dir):

tf.logging.error("Image directory '" + image_dir + "' not found.")

return None

end if

create an empty dictionary to store the results

result = {}

get a list of the sub-directories of the image directory

sub_dirs = [x[0] for x in gfile.Walk(image_dir)]

for each directory in the sub-directories list . . .

is_root_dir = True

for sub_dir in sub_dirs:

if we're on the 1st (root) directory, mark our boolean for that as false for the next time around and go back to the top of the for loop

if is_root_dir:

is_root_dir = False

continue

end if

dir_name = os.path.basename(sub_dir)

if dir_name == image_dir:

continue

end if

ToDo: This section should be refactored. The right way to do this would be to get a list of the files that are

ToDo: there then append (extend) those, not to get the name except the extension, then append an extension,

ToDo: this (current) way is error prone of the original file has an upper case or mixed case extension

extensions = ['.jpg', '.jpeg']

file_list = []

tf.logging.info("Looking for images in '" + dir_name + "'")

for extension in extensions:

file_glob = os.path.join(image_dir, dir_name, '*' + extension)

file_list.extend(gfile.Glob(file_glob))

end for

if the file list is empty at this point, log a warning and bail

```

if not file_list:

    tf.logging.warning('No files found')

    continue

# end if

# if the length of the file list is less than 20 or
more than the max number, log an applicable
warning (do not return, however)

if len(file_list) < 20:

    tf.logging.warning('WARNING: Folder has
less than 20 images, which may cause issues.')

    elif len(file_list) >
MAX_NUM_IMAGES_PER_CLASS:

        tf.logging.warning('WARNING: Folder {} has
more than {} images. Some images will never be
selected.'.format(dir_name,
MAX_NUM_IMAGES_PER_CLASS))

# end if

label_name = re.sub(r'^a-z0-9+', '',
dir_name.lower())

training_images = []

testing_images = []

validation_images = []

for file_name in file_list:

    base_name = os.path.basename(file_name)

    # We want to ignore anything after
'_nohash_' in the file name when deciding which
set to put an image in, the data set creator

    # has a way of grouping photos that are
close variations of each other. For example this is
used in the plant disease data set

    # to group multiple pictures of the same
leaf.

    hash_name = re.sub(r'_nohash_.*$', '',
file_name)

    # This looks a bit magical, but we need to
decide whether this file should go into the training,
testing, or validation sets,

    # and we want to keep existing files in the
same set even if more files are subsequently
added. To do that, we need a stable

```

```

# way of deciding based on just the file
name itself, so we do a hash of that and then use
that to generate a probability value

# that we use to assign it.

hash_name_hashed =
hashlib.sha1(compat.as_bytes(hash_name)).hexdigest()

percentage_hash =
((int(hash_name_hashed, 16) %
(MAX_NUM_IMAGES_PER_CLASS + 1)) * (100.0 /
MAX_NUM_IMAGES_PER_CLASS))

if percentage_hash < validation_percentage:

    validation_images.append(base_name)

elif percentage_hash < (testing_percentage
+ validation_percentage):

    testing_images.append(base_name)

else:

    training_images.append(base_name)

    result[label_name] = {'dir': dir_name,
'training': training_images, 'testing':
testing_images, 'validation': validation_images,}

return result

def add_jpeg_decoding(input_width, input_height,
input_depth, input_mean, input_std):

    """

    Adds operations that perform JPEG decoding and
resizing to the graph..

    Args:

        input_width: Desired width of the image fed
into the recognizer graph.

        input_height: Desired width of the image fed
into the recognizer graph.

        input_depth: Desired channels of the image
fed into the recognizer graph.

        input_mean: Pixel value that should be zero in
the image for the graph.

        input_std: How much to divide the pixel values
by before recognition.

    Returns:

```

Tensors for the node to feed JPEG data into,
and the output of the preprocessing steps.

```
"""

jpeg_data = tf.placeholder(tf.string,
name='DecodeJPGInput')

decoded_image =
tf.image.decode_jpeg(jpeg_data,
channels=input_depth)

decoded_image_as_float =
tf.cast(decoded_image, dtype=tf.float32)

decoded_image_4d =
tf.expand_dims(decoded_image_as_float, 0)

resize_shape = tf.stack([input_height,
input_width])

resize_shape_as_int = tf.cast(resize_shape,
dtype=tf.int32)

resized_image =
tf.image.resize_bilinear(decoded_image_4d,
resize_shape_as_int)

offset_image = tf.subtract(resized_image,
input_mean)

mul_image = tf.multiply(offset_image, 1.0 /
input_std)

return jpeg_data, mul_image

def add_input_distortions(flip_left_right,
random_crop, random_scale, random_brightness,
input_width, input_height,

input_depth, input_mean,
input_std):

jpeg_data = tf.placeholder(tf.string,
name='DistortJPGInput')

decoded_image =
tf.image.decode_jpeg(jpeg_data,
channels=input_depth)

decoded_image_as_float =
tf.cast(decoded_image, dtype=tf.float32)

decoded_image_4d =
tf.expand_dims(decoded_image_as_float, 0)

margin_scale = 1.0 + (random_crop / 100.0)

resize_scale = 1.0 + (random_scale / 100.0)
```

```
margin_scale_value = tf.constant(margin_scale)

resize_scale_value =
tf.random_uniform(tensor_shape.scalar(),
minval=1.0, maxval=resize_scale)

scale_value = tf.multiply(margin_scale_value,
resize_scale_value)

precrop_width = tf.multiply(scale_value,
input_width)

precrop_height = tf.multiply(scale_value,
input_height)

precrop_shape = tf.stack([precrop_height,
precrop_width])

precrop_shape_as_int = tf.cast(precrop_shape,
dtype=tf.int32)

precropped_image =
tf.image.resize_bilinear(decoded_image_4d,
precrop_shape_as_int)

precropped_image_3d =
tf.squeeze(precropped_image, squeeze_dims=[0])

cropped_image =
tf.random_crop(precropped_image_3d,
[input_height, input_width, input_depth])

if flip_left_right:

flipped_image =
tf.image.random_flip_left_right(cropped_image)

else:

flipped_image = cropped_image

# end if

brightness_min = 1.0 - (random_brightness /
100.0)

brightness_max = 1.0 + (random_brightness /
100.0)

brightness_value =
tf.random_uniform(tensor_shape.scalar(),
minval=brightness_min, maxval=brightness_max)

brightened_image = tf.multiply(flipped_image,
brightness_value)

offset_image = tf.subtract(brightened_image,
input_mean)

mul_image = tf.multiply(offset_image, 1.0 /
input_std)

distort_result = tf.expand_dims(mul_image, 0,
name='DistortResult')

return jpeg_data, distort_result
```

```
def cache_bottlenecks(sess, image_lists,
                      image_dir, bottleneck_dir, jpeg_data_tensor,
                      decoded_image_tensor,
```

```
                      resized_input_tensor,
                      bottleneck_tensor, architecture):
```

```
    """
```

Ensures all the training, testing, and validation
bottlenecks are cached.

Because we're likely to read the same image
multiple times (if there are no distortions applied
during training) it

can speed things up a lot if we calculate the
bottleneck layer values once for each image during
preprocessing,

and then just read those cached values
repeatedly during training. Here we go through all
the images we've found,

calculate those values, and save them off.

Args:

sess: The current active TensorFlow Session.

image_lists: Dictionary of training images for
each label.

image_dir: Root folder string of the subfolders
containing the training images.

bottleneck_dir: Folder string holding cached
files of bottleneck values.

jpeg_data_tensor: Input tensor for jpeg data
from file.

decoded_image_tensor: The output of
decoding and resizing the image.

resized_input_tensor: The input node of the
recognition graph.

bottleneck_tensor: The penultimate output
layer of the graph.

architecture: The name of the model
architecture.

Returns:

Nothing.

```
    """
```

```
    how_many_bottlenecks = 0
```

```
    makeDirIfDoesNotExist(bottleneck_dir)
```

```
    for label_name, label_lists in image_lists.items():
```

```
        for category in ['training', 'testing',
                        'validation']:
```

```
            category_list = label_lists[category]
```

```
            for index, unused_base_name in
                enumerate(category_list):
```

```
                get_or_create_bottleneck(sess,
                    image_lists, label_name, index, image_dir,
                    category, bottleneck_dir,
```

```
                    jpeg_data_tensor,
                    decoded_image_tensor, resized_input_tensor,
                    bottleneck_tensor, architecture)
```

```
            # end for
```

```
            how_many_bottlenecks += 1
```

```
            if how_many_bottlenecks % 100 == 0:
```

```
                tf.logging.info(str(how_many_bottlenecks) + '
                    bottleneck files created.')
```

```
            # end if
```

```
def get_or_create_bottleneck(sess, image_lists,
                              label_name, index, image_dir, category,
                              bottleneck_dir, jpeg_data_tensor,
```

```
                              decoded_image_tensor,
                              resized_input_tensor, bottleneck_tensor,
                              architecture):
```

```
    """
```

Retrieves or calculates bottleneck values for an
image.

If a cached version of the bottleneck data exists
on-disk, return that, otherwise calculate the data
and save it to disk for future use.

Args:

sess: The current active TensorFlow Session.

image_lists: Dictionary of training images for
each label.

label_name: Label string we want to get an image for.

index: Integer offset of the image we want. This will be modulo-ed by the available number of images for the label, so it can be arbitrarily large.

image_dir: Root folder string of the subfolders containing the training images.

category: Name string of which set to pull images from - training, testing, or validation.

bottleneck_dir: Folder string holding cached files of bottleneck values.

jpeg_data_tensor: The tensor to feed loaded jpeg data into.

decoded_image_tensor: The output of decoding and resizing the image.

resized_input_tensor: The input node of the recognition graph.

bottleneck_tensor: The output tensor for the bottleneck values.

architecture: The name of the model architecture.

Returns:

Numpy array of values produced by the bottleneck layer for the image.

```
"""
```

```
label_lists = image_lists[label_name]
```

```
sub_dir = label_lists['dir']
```

```
sub_dir_path = os.path.join(bottleneck_dir,
sub_dir)
```

```
makeDirIfDoesNotExist(sub_dir_path)
```

```
bottleneck_path =
get_bottleneck_path(image_lists, label_name,
index, bottleneck_dir, category, architecture)
```

```
if not os.path.exists(bottleneck_path):
```

```
    create_bottleneck_file(bottleneck_path,
image_lists, label_name, index, image_dir,
category, sess, jpeg_data_tensor,
```

```
                           decoded_image_tensor,
resized_input_tensor, bottleneck_tensor)
```

```
# end if
```

```
# read in the contents of the bottleneck file as
one big string
```

```
with open(bottleneck_path, 'r') as
bottleneck_file:
```

```
    bottleneckBigString = bottleneck_file.read()
```

```
# end with
```

```
bottleneckValues = []
```

```
errorOccurred = False
```

```
try:
```

```
    # split the bottleneck file contents read in as
one big string into individual float values
```

```
    bottleneckValues = [float(individualString) for
individualString in bottleneckBigString.split(',')]
```

```
except ValueError:
```

```
    tf.logging.warning('Invalid float found,
recreating bottleneck')
```

```
    errorOccurred = True
```

```
# end try
```

```
if errorOccurred:
```

```
    # if an error occurred above, create (or re-
create) the bottleneck file
```

```
    create_bottleneck_file(bottleneck_path,
image_lists, label_name, index, image_dir,
category, sess,
```

```
                           jpeg_data_tensor,
decoded_image_tensor, resized_input_tensor,
bottleneck_tensor)
```

```
# read in the contents of the newly created
bottleneck file
```

```
with open(bottleneck_path, 'r') as
bottleneck_file:
```

```
    bottleneckBigString = bottleneck_file.read()
```

```
# end with
```

```
# split the bottleneck file contents read in as
one big string into individual float values again
```

```

        bottleneckValues = [float(individualString) for
individualString in bottleneckBigString.split(',')]

```

```

# end if

```

```

return bottleneckValues

```

```

def get_bottleneck_path(image_lists, label_name,
index, bottleneck_dir, category, architecture):

```

```

    """

```

Returns a path to a bottleneck file for a label at the given index.

Args:

image_lists: Dictionary of training images for each label.

label_name: Label string we want to get an image for.

index: Integer offset of the image we want. This will be moduloed by the

available number of images for the label, so it can be arbitrarily large.

bottleneck_dir: Folder string holding cached files of bottleneck values.

category: Name string of set to pull images from - training, testing, or

validation.

architecture: The name of the model architecture.

Returns:

File system path string to an image that meets the requested parameters.

```

    """

```

```

    return get_image_path(image_lists, label_name,
index, bottleneck_dir, category) + '_' + architecture
+ '.txt'

```

```

def create_bottleneck_file(bottleneck_path,
image_lists, label_name, index,

```

```

    image_dir, category, sess,
jpeg_data_tensor,

```

```

    decoded_image_tensor,
resized_input_tensor,

```

```

    bottleneck_tensor):

```

```

    """Create a single bottleneck file."""

```

```

    tf.logging.info('Creating bottleneck at ' +
bottleneck_path)

```

```

    image_path = get_image_path(image_lists,
label_name, index, image_dir, category)

```

```

    if not gfile.Exists(image_path):

```

```

        tf.logging.fatal('File does not exist %s',
image_path)

```

```

    # end if

```

```

    image_data = gfile.GFile(image_path,
'rb').read()

```

```

    try:

```

```

        bottleneck_values =
run_bottleneck_on_image(sess, image_data,
jpeg_data_tensor, decoded_image_tensor,
resized_input_tensor, bottleneck_tensor)

```

```

    except Exception as e:

```

```

        raise RuntimeError('Error during processing
file %s (%s)' % (image_path, str(e)))

```

```

    # end try

```

```

    bottleneck_string = ','.join(str(x) for x in
bottleneck_values)

```

```

    with open(bottleneck_path, 'w') as
bottleneck_file:

```

```

        bottleneck_file.write(bottleneck_string)

```

```

def run_bottleneck_on_image(sess, image_data,
image_data_tensor, decoded_image_tensor,
resized_input_tensor, bottleneck_tensor):

```

```

    """

```

Runs inference on an image to extract the 'bottleneck' summary layer.

Args:

sess: Current active TensorFlow Session.

image_data: String of raw JPEG data.

image_data_tensor: Input data layer in the graph.

decoded_image_tensor: Output of initial image resizing and preprocessing.

resized_input_tensor: The input node of the recognition graph.

bottleneck_tensor: Layer before the final softmax.

Returns:

Numpy array of bottleneck values.

"""

First decode the JPEG image, resize it, and rescale the pixel values.

```
resized_input_values =  
sess.run(decoded_image_tensor,  
{image_data_tensor: image_data})
```

Then run it through the recognition network.

```
bottleneck_values = sess.run(bottleneck_tensor,  
{resized_input_tensor: resized_input_values})
```

```
bottleneck_values =  
np.squeeze(bottleneck_values)
```

return bottleneck_values

def get_image_path(image_lists, label_name,
index, image_dir, category):

"""

Returns a path to an image for a label at the given index.

Args:

image_lists: Dictionary of training images for each label.

label_name: Label string we want to get an image for.

index: Int offset of the image we want. This will be modulated by the available number of images for the label, so it can be arbitrarily large.

image_dir: Root folder string of the subfolders containing the training images.

category: Name string of set to pull images from - training, testing, or validation.

Returns:

File system path string to an image that meets the requested parameters.

"""

if label_name not in image_lists:

```
tf.logging.fatal('Label does not exist %s.',  
label_name)
```

end if

label_lists = image_lists[label_name]

if category not in label_lists:

```
tf.logging.fatal('Category does not exist %s.',  
category)
```

end if

category_list = label_lists[category]

if not category_list:

```
tf.logging.fatal('Label %s has no images in the  
category %s.', label_name, category)
```

end if

mod_index = index % len(category_list)

base_name = category_list[mod_index]

sub_dir = label_lists['dir']

```
full_path = os.path.join(image_dir, sub_dir,  
base_name)
```

return full_path

```
def add_final_training_ops(class_count,  
final_tensor_name, bottleneck_tensor,  
bottleneck_tensor_size, quantize_layer):
```

"""

Adds a new softmax and fully-connected layer for training.

We need to retrain the top layer to identify our new classes, so this function

adds the right operations to the graph, along with some variables to hold the

weights, and then sets up all the gradients for the backward pass.

The set up for the softmax and fully-connected layers is based on:

<https://www.tensorflow.org/versions/master/tutorials/mnist/beginners/index.html>

Args:

class_count: Integer of how many categories of things we're trying to recognize.

final_tensor_name: Name string for the new final node that produces results.

bottleneck_tensor: The output of the main CNN graph.

bottleneck_tensor_size: How many entries in the bottleneck vector.

quantize_layer: Boolean, specifying whether the newly added layer should be quantized.

Returns:

The tensors for the training and cross entropy results, and tensors for the bottleneck input and ground truth input.

```
"""

with tf.name_scope('input'):

    bottleneck_input =
    tf.placeholder_with_default(bottleneck_tensor,
    shape=[None, bottleneck_tensor_size],
    name='BottleneckInputPlaceholder')

    ground_truth_input = tf.placeholder(tf.int64,
    [None], name='GroundTruthInput')

# end with

# Organizing the following ops as
`final_training_ops` so they're easier to see in
TensorBoard

layer_name = 'final_training_ops'

with tf.name_scope(layer_name):

    quantized_layer_weights = None

    quantized_layer_biases = None
```

```
with tf.name_scope('weights'):

    initial_value =
    tf.truncated_normal([bottleneck_tensor_size,
    class_count], stddev=0.001)

    layer_weights = tf.Variable(initial_value,
    name='final_weights')

    if quantize_layer:

        quantized_layer_weights =
        quant_ops.MovingAvgQuantize(layer_weights,
        is_training=True)

attachTensorBoardSummaries(quantized_layer_weights)

# end if

# this comment is necessary to suppress an
unnecessary PyCharm warning

# noinspection PyTypeChecker

attachTensorBoardSummaries(layer_weights)

# end with

with tf.name_scope('biases'):

    layer_biases =
    tf.Variable(tf.zeros([class_count]),
    name='final_biases')

    if quantize_layer:

        quantized_layer_biases =
        quant_ops.MovingAvgQuantize(layer_biases,
        is_training=True)

attachTensorBoardSummaries(quantized_layer_biases)

# end if

# this comment is necessary to suppress an
unnecessary PyCharm warning

# noinspection PyTypeChecker

attachTensorBoardSummaries(layer_biases)

# end with

with tf.name_scope('Wx_plus_b'):

    if quantize_layer:
```

```

        logits = tf.matmul(bottleneck_input,
                           quantized_layer_weights) +
                           quantized_layer_biases

        logits =
        quant_ops.MovingAvgQuantize(logits, init_min=-
        32.0, init_max=32.0, is_training=True, num_bits=8,

                                   narrow_range=False,
        ema_decay=0.5)

        tf.summary.histogram('pre_activations',
        logits)

        else:

            logits = tf.matmul(bottleneck_input,
            layer_weights) + layer_biases

            tf.summary.histogram('pre_activations',
            logits)

        final_tensor = tf.nn.softmax(logits,
        name=final_tensor_name)

        tf.summary.histogram('activations', final_tensor)

        with tf.name_scope('cross_entropy'):

            cross_entropy_mean =
            tf.losses.sparse_softmax_cross_entropy(labels=ground_truth_input, logits=logits)

            tf.summary.scalar('cross_entropy',
            cross_entropy_mean)

        with tf.name_scope('train'):

            optimizer =
            tf.train.GradientDescentOptimizer(LEARNING_RATE)

            train_step =
            optimizer.minimize(cross_entropy_mean)

            return (train_step, cross_entropy_mean,
            bottleneck_input, ground_truth_input,
            final_tensor)

def attachTensorBoardSummaries(var):

```

```

        """Attach a lot of summaries to a Tensor (for
        TensorBoard visualization)."""

        with tf.name_scope('summaries'):

            mean = tf.reduce_mean(var)

            tf.summary.scalar('mean', mean)

            with tf.name_scope('stddev'):

                stddev =
                tf.sqrt(tf.reduce_mean(tf.square(var - mean)))

            # end with

            tf.summary.scalar('stddev', stddev)

            tf.summary.scalar('max', tf.reduce_max(var))

            tf.summary.scalar('min', tf.reduce_min(var))

            tf.summary.histogram('histogram', var)

def add_evaluation_step(result_tensor,
        ground_truth_tensor):

    """

    Inserts the operations we need to evaluate the
    accuracy of our results.

    Args:

        result_tensor: The new final node that
        produces results.

        ground_truth_tensor: The node we feed
        ground truth data into.

    Returns:

        Tuple of (evaluation step, prediction).

    """

    with tf.name_scope('accuracy'):

        with tf.name_scope('correct_prediction'):

            prediction = tf.argmax(result_tensor, 1)

            correct_prediction = tf.equal(prediction,
            ground_truth_tensor)

            # end with

            with tf.name_scope('accuracy'):

```

```

        evaluation_step =
tf.reduce_mean(tf.cast(correct_prediction,
tf.float32))

    # end with

    tf.summary.scalar('accuracy', evaluation_step)

    return evaluation_step, prediction

```

```

def get_random_distorted_bottlenecks(sess,
image_lists, how_many, category, image_dir,
input_jpeg_tensor, distorted_image,

```

```

        resized_input_tensor,
bottleneck_tensor):

```

```

    """

```

Retrieves bottleneck values for training images, after distortions.

If we're training with distortions like crops, scales, or flips, we have to recalculate the full model for every image,

and so we can't use cached bottleneck values. Instead we find random images for the requested category, run them through

the distortion graph, and then the full graph to get the bottleneck results for each.

Args:

sess: Current TensorFlow Session.

image_lists: Dictionary of training images for each label.

how_many: The integer number of bottleneck values to return.

category: Name string of which set of images to fetch - training, testing, or validation.

image_dir: Root folder string of the subfolders containing the training images.

input_jpeg_tensor: The input layer we feed the image data to.

distorted_image: The output node of the distortion graph.

resized_input_tensor: The input node of the recognition graph.

bottleneck_tensor: The bottleneck output layer of the CNN graph.

Returns:

List of bottleneck arrays and their corresponding ground truths.

```

    """

```

```

    class_count = len(image_lists.keys())

```

```

    bottlenecks = []

```

```

    ground_truths = []

```

```

    for unused_i in range(how_many):

```

```

        label_index = random.randrange(class_count)

```

```

        label_name =
list(image_lists.keys())[label_index]

```

```

        image_index =
random.randrange(MAX_NUM_IMAGES_PER_CLASS + 1)

```

```

        image_path = get_image_path(image_lists,
label_name, image_index, image_dir, category)

```

```

        if not gfile.Exists(image_path):

```

```

            tf.logging.fatal('File does not exist %s',
image_path)

```

```

        # end if

```

```

        jpeg_data = gfile.GFile(image_path,
'rb').read()

```

```

        # Note that we materialize the
distorted_image_data as a numpy array before

```

```

        # sending running inference on the image. This
involves 2 memory copies and

```

```

        # might be optimized in other
implementations.

```

```

        distorted_image_data =
sess.run(distorted_image, {input_jpeg_tensor:
jpeg_data})

```

```

        bottleneck_values =
sess.run(bottleneck_tensor, {resized_input_tensor:
distorted_image_data})

```

```

        bottleneck_values =
np.squeeze(bottleneck_values)

```

```

        bottlenecks.append(bottleneck_values)

```

```

        ground_truths.append(label_index)

```

```

    return bottlenecks, ground_truths

```

```
def get_random_cached_bottlenecks(sess,
image_lists, how_many, category, bottleneck_dir,
image_dir, jpeg_data_tensor,
```

```
    decoded_image_tensor,
resized_input_tensor, bottleneck_tensor,
architecture):
```

```
    """
```

Retrieves bottleneck values for cached images.

If no distortions are being applied, this function can retrieve the cached bottleneck values directly from disk for

images. It picks a random set of images from the specified category.

Args:

sess: Current TensorFlow Session.

image_lists: Dictionary of training images for each label.

how_many: If positive, a random sample of this size will be chosen. If negative, all bottlenecks will be retrieved.

category: Name string of which set to pull from - training, testing, or validation.

bottleneck_dir: Folder string holding cached files of bottleneck values.

image_dir: Root folder string of the subfolders containing the training images.

jpeg_data_tensor: The layer to feed jpeg image data into.

decoded_image_tensor: The output of decoding and resizing the image.

resized_input_tensor: The input node of the recognition graph.

bottleneck_tensor: The bottleneck output layer of the CNN graph.

architecture: The name of the model architecture.

Returns:

List of bottleneck arrays, their corresponding ground truths, and the relevant filenames.

```
    """
```

```
    class_count = len(image_lists.keys())
```

```
    bottlenecks = []
```

```
    ground_truths = []
```

```
    filenames = []
```

```
    if how_many >= 0:
```

```
        # Retrieve a random sample of bottlenecks.
```

```
        for unused_i in range(how_many):
```

```
            label_index =
random.randrange(class_count)
```

```
            label_name =
list(image_lists.keys())[label_index]
```

```
            image_index =
random.randrange(MAX_NUM_IMAGES_PER_CLASS + 1)
```

```
            image_name = get_image_path(image_lists,
label_name, image_index, image_dir, category)
```

```
            bottleneck = get_or_create_bottleneck(sess,
image_lists, label_name, image_index, image_dir,
category, bottleneck_dir,
```

```
                jpeg_data_tensor,
decoded_image_tensor, resized_input_tensor,
bottleneck_tensor, architecture)
```

```
            bottlenecks.append(bottleneck)
```

```
            ground_truths.append(label_index)
```

```
            filenames.append(image_name)
```

```
        # end for
```

```
    else:
```

```
        # Retrieve all bottlenecks.
```

```
        for label_index, label_name in
enumerate(image_lists.keys()):
```

```
            for image_index, image_name in
enumerate(image_lists[label_name][category]):
```

```
                image_name =
get_image_path(image_lists, label_name,
image_index, image_dir, category)
```

```
                bottleneck =
get_or_create_bottleneck(sess, image_lists,
label_name, image_index, image_dir, category,
bottleneck_dir,
```

```

        jpeg_data_tensor,
        decoded_image_tensor, resized_input_tensor,
        bottleneck_tensor, architecture)

```

```

        bottlenecks.append(bottleneck)

```

```

        ground_truths.append(label_index)

```

```

        filenames.append(image_name)

```

```

    return bottlenecks, ground_truths, filenames

```

```

def save_graph_to_file(sess, graph,
    graph_file_name):

```

```

    output_graph_def =
    graph_util.convert_variables_to_constants(sess,
    graph.as_graph_def(), [FINAL_TENSOR_NAME])

```

```

    with gfile.GFile(graph_file_name, 'wb') as f:

```

```

        f.write(output_graph_def.SerializeToString())

```

```

    # end with

```

```

    return

```

```

if __name__ == '__main__':

```

```

    main()

```

En este otro código es donde se prueba la base de datos usando imágenes fuera de las carpetas ya estudiadas por el programa. Al igual que el código anterior esta marcado con comentarios para su mejor comprensión y entendimiento.

```

# test.py

```

```

import os

```

```

import tensorflow as tf

```

```

import numpy as np

```

```

import cv2

```

```

RETRAINED_LABELS_TXT_FILE_LOC = os.getcwd() +
"/" + "retrained_labels.txt"

```

```

RETRAINED_GRAPH_PB_FILE_LOC = os.getcwd() +
"/" + "retrained_graph.pb"

```

```

TEST_IMAGES_DIR = os.getcwd() + "/test_images"

```

```

SCALAR_RED = (0.0, 0.0, 255.0)

```

```

SCALAR_BLUE = (255.0, 0.0, 0.0)

```

```

def main():

```

```

    print("starting program . . .")

```

```

    if not checkIfNecessaryPathsAndFilesExist():

```

```

        return

```

```

# end if

```

```

# get a list of classifications from the labels file

```

```

classifications = []

```

```

# for each line in the label file . . .

```

```

    for currentLine in
    tf.gfile.GFile(RETRAINED_LABELS_TXT_FILE_LOC):

```

```

        # remove the carriage return

```

```

        classification = currentLine.rstrip()

```

```

        # and append to the list

```

```

        classifications.append(classification)

```

```

    # end for

```

```

# show the classifications to prove out that we
were able to read the label file successfully

```

```

    print("classifications = " + str(classifications))

```

```

# load the graph from file

```

```

    with
    tf.gfile.GFile(RETRAINED_GRAPH_PB_FILE_LOC
    , 'rb') as retrainedGraphFile:

```

```

        # instantiate a GraphDef object

```

```

        graphDef = tf.GraphDef()

```

```

        # read in retrained graph into the GraphDef
        object

graphDef.ParseFromString(retrainedGraphFile.read
())

        # import the graph into the current default
        Graph, note that we don't need to be concerned
        with the return value

        _ = tf.import_graph_def(graphDef, name="")

        # end with

        # if the test image directory listed above is not
        valid, show an error message and bail

        if not os.path.isdir(TEST_IMAGES_DIR):

            print("the test image directory does not seem
            to be a valid directory, check file / directory paths")

            return

        # end if

        with tf.Session() as sess:

            # for each file in the test images directory . . .

            for fileName in os.listdir(TEST_IMAGES_DIR):

                # if the file does not end in .jpg or .jpeg
                (case-insensitive), continue with the next iteration
                of the for loop

                if not (fileName.lower().endswith(".jpg") or
                fileName.lower().endswith(".jpeg")):

                    continue

                # end if

                # show the file name on std out

                print(fileName)

                # get the file name and full path of the
                current image file

                imageFilePath =
                os.path.join(TEST_IMAGES_DIR, fileName)

                # attempt to open the image with OpenCV

```

```

        openCVImage =
        cv2.imread(imageFilePath)

        # if we were not able to successfully open
        the image, continue with the next iteration of the
        for loop

        if openCVImage is None:

            print("unable to open " + fileName + " as
            an OpenCV image")

            continue

        # end if

        # get the final tensor from the graph

        finalTensor =
        sess.graph.get_tensor_by_name('final_result:0')

        # convert the OpenCV image (numpy array)
        to a TensorFlow image

        tfImage = np.array(openCVImage)[:,:,:,:3]

        # run the network to get the predictions

        predictions = sess.run(finalTensor,
        {'DecodeJpeg:0': tfImage})

        # sort predictions from most confidence to
        least confidence

        sortedPredictions = predictions[0].argsort()[
        -len(predictions[0]):::-1]

        print("-----")

        # keep track of if we're going through the
        next for loop for the first time so we can show
        more info about

        # the first prediction, which is the most
        likely prediction (they were sorted descending
        above)

        onMostLikelyPrediction = True

        # for each prediction . . .

```

```

        for prediction in sortedPredictions:

            strClassification =
            classifications[prediction]

            # if the classification (obtained from the
            directory name) ends with the letter "s", remove
            the "s" to change from plural to singular

            if strClassification.endswith("s"):

                strClassification = strClassification[:-1]

            # end if

            # get confidence, then get confidence
            rounded to 2 places after the decimal

            confidence = predictions[0][prediction]

            # if we're on the first (most likely)
            prediction, state what the object appears to be and
            show a % confidence to two decimal places

            if onMostLikelyPrediction:

                # get the score as a %

                scoreAsAPercent = confidence * 100.0

                # show the result to std out

                print("the object appears to be a " +
                strClassification + ", " +
                "{0:.2f}".format(scoreAsAPercent) + "%
                confidence")

                # write the result on the image

                writeResultOnImage(openCVImage,
                strClassification + ", " +
                "{0:.2f}".format(scoreAsAPercent) + "%
                confidence")

            # finally we can show the OpenCV
            image

            cv2.imshow(fileName, openCVImage)

            print("Guardar Imagen?")

            # mark that we've show the most likely
            prediction at this point so the additional
            information in

            # this if statement does not show again
            for this image

            onMostLikelyPrediction = False

```

```

        # end if

        # for any prediction, show the confidence
        as a ratio to five decimal places

        print(strClassification + " (" +
        "{0:.5f}".format(confidence) + ")")

        # end for

        # pause until a key is pressed so the user can
        see the current image (shown above) and the
        prediction info

        cv2.waitKey()

        # after a key is pressed, close the current
        window to prep for the next time around

        cv2.destroyAllWindows()

        # end for

        # end with

        # write the graph to file so we can view with
        TensorBoard

        tfFileWriter =
        tf.summary.FileWriter(os.getcwd())

        tfFileWriter.add_graph(sess.graph)

        tfFileWriter.close()

def checkIfNecessaryPathsAndFilesExist():

    if not os.path.exists(TEST_IMAGES_DIR):

        print("")

        print('ERROR: TEST_IMAGES_DIR "' +
        TEST_IMAGES_DIR + '" does not seem to exist')

        print('Did you set up the test images?')

        print("")

        return False

    # end if

```

```

    if not
os.path.exists(RETRAINED_LABELS_TXT_FILE_LOC):

        print('ERROR:
RETRAINED_LABELS_TXT_FILE_LOC "' +
RETRAINED_LABELS_TXT_FILE_LOC + '" does not
seem to exist')

        return False

    # end if

    if not
os.path.exists(RETRAINED_GRAPH_PB_FILE_LOC):

        print('ERROR:
RETRAINED_GRAPH_PB_FILE_LOC "' +
RETRAINED_GRAPH_PB_FILE_LOC + '" does not
seem to exist')

        return False

    # end if

    return True

def writeResultOnImage(openCVImage,
resultText):

    # ToDo: this function may take some further
fine-tuning to show the text well given any possible
image size

    imageHeight, imageWidth, sceneNumChannels =
openCVImage.shape

    # choose a font

    fontFace = cv2.FONT_HERSHEY_TRIPLEX

    # chose the font size and thickness as a fraction
of the image size

    fontScale = 1.0

    fontThickness = 2

    # make sure font thickness is an integer, if not,
the OpenCV functions that use this may crash

    fontThickness = int(fontThickness)

    upperLeftTextOriginX = int(imageWidth * 0.05)

    upperLeftTextOriginY = int(imageHeight * 0.05)

    textSize, baseline = cv2.getTextSize(resultText,
fontFace, fontScale, fontThickness)

    textSizeWidth, textSizeHeight = textSize

    # calculate the lower left origin of the text area
based on the text area center, width, and height

    lowerLeftTextOriginX = upperLeftTextOriginX

    lowerLeftTextOriginY = upperLeftTextOriginY +
textSizeHeight

    # write the text on the image

    cv2.putText(openCVImage, resultText,
(lowerLeftTextOriginX, lowerLeftTextOriginY),
fontFace, fontScale, SCALAR_BLUE, fontThickness)

if __name__ == "__main__":

    main()

```

Este proyecto si se implementa con tanta información e imágenes puede llegar a distinguir casi cualquier cosa, al igual que se puede hacer en un app para usar el programa a tiempo real, es por eso de su amplio espectro de escalabilidad. Lo bueno de este proyecto es que solo necesita dos códigos, pero se puede convertir en uno si el usuario sabe hacer todo el proceso antes. El programa se puede personalizar a gustos de los usuarios al igual que este enfocado a un objetivo más directo. Por ultimo este programa puede tener una gran seguridad de usuario al igual que ser anónimo en su uso para la discreción de los clientes.

En conclusión, el proyecto para la clasificación de imágenes es muy simple de usar y efectivo, la única que si se necesita es el tiempo de recolección de imágenes, tener la paciencia para que el programa aprende a diferenciarlos y al probarlo.

Referencias:

<https://www.apsl.net/blog/2017/12/05/tensor-flow-para-principiantes-i/>

https://www.youtube.com/watch?v=szNPBn_RBfA

<https://www.youtube.com/watch?v=90gpNF3KzK8&t=72s>